

CAPSTONE PROJECT 7 PRODUCTIONIZATION OF ML SYSTEMS DOCUMENTATION

This document explains the complete MLOps workflow for the Travel Price Prediction capstone project. It covers all seven stages: REST API, Streamlit App, Docker Deployment, Kubernetes Deployment, Airflow Scheduling, CI/CD Pipeline, and MLflow Tracking. Each section describes the workflow and references the screenshots taken during implementation.

STAGE 1: REST API Flask Deployment

WORKFLOW OVERVIEW:

The trained machine learning models (flight price regression and gender classification) are served as a REST API using Flask. The API exposes two prediction endpoints `/predict_price` and `/predict_gender` that accept POST requests with flight/user features in JSON format and return real time predictions.

HOW IT WORKS:


1. The Flask app (`app.py`) loads the saved joblib model files at startup.
2. When a POST request hits `/predict_price`, the input features are extracted, scaled using the saved `StandardScaler`, and passed to the `GradientBoostingRegressor` model.
3. The predicted price is returned as a JSON response.
4. Similarly, `/predict_gender` uses the classification model to return the predicted gender label.
5. A root GET endpoint (`/`) returns a health check message confirming the API is live.

SCREENSHOT:

```
(base) karthikdoguparthi@Lazys-Mac Module 7 % source .venv/bin/activate
(.venv) (base) karthikdoguparthi@Lazys-Mac Module 7 % python app.py
* Serving Flask app 'app'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5001
* Running on http://192.168.29.35:5001
Press CTRL-C to quit
192.168.29.35 - - [06/May/2026 17:28:20] "GET / HTTP/1.1" 200 -
192.168.29.35 - - [06/May/2026 17:28:20] "GET /favicon.ico HTTP/1.1" 404 -
127.0.0.1 - - [06/May/2026 17:28:27] "GET / HTTP/1.1" 200 -
127.0.0.1 - - [06/May/2026 17:34:01] "POST /predict_gender HTTP/1.1" 200 -
127.0.0.1 - - [06/May/2026 17:34:33] "POST /predict_gender HTTP/1.1" 200 -
127.0.0.1 - - [06/May/2026 17:34:43] "POST /predict_price HTTP/1.1" 200 -
127.0.0.1 - - [06/May/2026 17:41:47] "POST /predict_price HTTP/1.1" 200 -
127.0.0.1 - - [06/May/2026 17:42:13] "POST /predict_gender HTTP/1.1" 200 -
^
```

Terminal showing Flask server running

The terminal shows the Flask development server running with the command `"python app.py"` inside a virtual environment (`.venv`). The server is active on all addresses (`0.0.0.0`), accessible at both `http://127.0.0.1:5001` and `http://192.168.29.35:5001`. The server logs confirm successful POST requests to `/predict_gender` and `/predict_price` endpoints, all returning HTTP 200 status codes confirming the API is fully operational.



Browser showing API health check

Navigating to `http://192.168.29.35:5001` in the browser returns the JSON response: `{ "message": "Travel Prediction API is running!" }` confirming the Flask API root endpoint is live and accessible on the local network.

KEY ENDPOINTS:

- GET / – Health check
- POST /predict_price – Returns predicted flight price
- POST /predict_gender – Returns predicted user gender

STAGE 2: STREAMLIT APP


WORKFLOW OVERVIEW:

A Streamlit web application provides an interactive dashboard for the Travel Price Prediction system. It is split into three sections: EDA & Insights, Hotel Recommendations, and Model Predictions. The app communicates with the Flask API for real-time predictions.

HOW IT WORKS:


1. On launch, the app loads the three datasets (flights.csv, hotels.csv, users.csv) and displays a "Data loaded successfully!" confirmation toast.
2. The sidebar dropdown allows the user to switch between three sections.
3. The EDA & Insights section displays visualisations including a boxplot of flight price by flight type and a histogram of user age distribution.
4. The Hotel Recommendations section accepts a User Code input and returns personalised hotel suggestions using the collaborative filtering model.
5. The Model Predictions section connects to the running Flask API at `http://localhost:5001` and submits flight or user details to get real-time predictions.

SCREENSHOT:




EDA & Insights section

The Streamlit app is running at localhost:8501. The "EDA & Insights" section is selected from the sidebar. Two charts are visible: a boxplot showing Flight Price Distribution by Flight Type (firstClass prices are significantly higher than economic and premium), and a histogram of User Age Distribution with a KDE overlay. The "Data loaded successfully!" toast confirms all three CSV files were loaded correctly.




Hotel Recommendations section

The "Hotel Recommendations" section is selected. The user enters User Code = 1 and clicks "Get Recommendations". The system returns a ranked table of personalised hotel suggestions: Hotel BW (6 bookings), Hotel BD (3), Hotel Z (3), Hotel A (2), Hotel K (2). These are derived from cosine similarity between user booking histories.



Model Predictions: Flight Price

The "Model Predictions" section is selected with the "Flight Price" task chosen. Input values are: From = Recife (PE), To = Florianopolis (SC), Flight Type = firstClass, Agency = FlyingDrops, Time = 2.00 hours, Distance = 500 km, Date = 2026/05/06, Price per KM = 2.00, Time Dist Ratio = 0.00. After clicking "Predict Price", the app calls the Flask API and displays: Predicted Price: \$1224.30



Model Predictions: User Gender

The same Model Predictions section is shown with the "User Gender" task selected. Input values include Flight Type = firstClass, Agency = FlyingDrops, Time = 2.00 hours, Distance = 500 km, Date = 2026/05/06, Price Paid = 1000.00. After clicking "Predict Gender", the API returns: Predicted Gender: none

STAGE 3: DOCKER DEPLOYMENT

WORKFLOW OVERVIEW:

The Flask API is containerised using Docker, making it environment-independent and portable. The Dockerfile packages the application code, model files, and all dependencies into a single Docker image. The container is then run and exposed on port 5001.

HOW IT WORKS:

1. The Dockerfile uses python:3.9 slim as the base image to keep the container lightweight.
2. The working directory is set to /app, all project files are copied in, and dependencies are installed from requirements.txt.
3. The image is built with the tag "travel-api".
4. The container is launched with port 5001 mapped to the host and named "travel-container".
5. Docker Desktop is used to monitor the running container's health, CPU, and memory usage.

SCREENSHOT:

```
(base) karthikdoguparthi@Lazys-Mac Module 7 % docker build -t travel-api .
docker run -d -p 5001:5001 --name travel-container travel-api

[+] Building 131.9s (10/10) FINISHED                                docker:desktop-linux
=> [internal] load build definition from Dockerfile                  0.0s
=> => transferring dockerfile: 555B                                  0.0s
=> [internal] load metadata for docker.io/library/python:3.9-slim   5.6s
=> [auth] library/python:pull token for registry-1.docker.io        0.0s
=> [internal] load .dockerignore                                     0.0s
=> => transferring context: 2B                                         0.0s
=> [1/4] FROM docker.io/library/python:3.9-slim@sha256:2d97f6910b16bd338d 6.3s
=> => resolve docker.io/library/python:3.9-slim@sha256:2d97f6910b16bd338d 0.0s
=> => sha256:c23f4b50347300e01a1a1da6dd0266adc8e44671002ad28 251B / 251B 0.5s
=> => sha256:479b8ad8bcc3edbeba82d4959dfbdc65226d5b55df3f 1.27MB / 1.27MB 1.1s
=> => sha256:ebfe7d4fa0c501a81a5ba6d1e1e2958e4b005d3ce3 13.83MB / 13.83MB 2.7s
=> => sha256:a16e551192670581ec8359c70ab9eebf8f2af5468f 30.14MB / 30.14MB 4.5s
=> => extracting sha256:a16e551192670581ec8359c70ab9eebf8f2af5468f79b3d 1.0s
=> => extracting sha256:479b8ad8bcc3edbeba82d4959dfbdc65226d5b55df3f36914 0.1s
=> => extracting sha256:ebfe7d4fa0c501a81a5ba6d1e1e2958e4b005d3ce3827b0ad 0.8s
=> => extracting sha256:c23f4b50347300e01a1a1da6dd0266adc8e44671002ad28c 0.0s
=> [internal] load build context                                     9.6s
=> => transferring context: 907.35MB                                   9.5s
=> [2/4] WORKDIR /app                                              0.7s
=> [3/4] COPY . /app                                              3.8s
=> [4/4] RUN pip install --no-cache-dir -r requirements.txt       82.5s
=> => exporting layers                                                30.2s
=> => exporting manifest sha256:f3568f3dadc8a91765d2a3c1fd2a244c82d585d55 20.1s
=> => exporting config sha256:87b8856a70ee8e0a0647d7dbdda0b65d19ac8242efb 0.0s
=> => exporting attestation manifest sha256:44f20cc31452e5a501dfe76a9e998 0.0s
=> => exporting manifest list sha256:40862001f9b7536fa022a4b7fd2dac3d9a63 0.0s
=> => naming to docker.io/library/travel-api:latest                  0.0s
=> => unpacking to docker.io/library/travel-api:latest              10.0s
e238385af19201746b4e54973d400fb172a81a3b5b4db7fc4fb10d20540c111e
(base) karthikdoguparthi@Lazys-Mac Module 7 %
```


Terminal: Docker build and run output

The terminal shows the complete output of:

```
docker build -t travel-api .
```

```
docker run -d -p 5001:5001 --name travel-container travel-api
```

The build completed successfully in 131.9 seconds across 10 build steps (10/10 FINISHED). Key steps include: pulling the python:3.9-slim base image, loading the Dockerfile, copying the 907MB build context (includes model files), installing all pip dependencies (82.5s), exporting layers, and naming the final image as docker.io/library/travel-api:latest. A container ID is returned confirming the container started.



Docker Desktop: Container running

Docker Desktop shows the "travel-container" is actively running (green dot) from the "travel-api" image. The container ID is e238385af192, port mapping is 5001:5001, CPU usage is 0.02%, and memory is 104.2MB out of 7.57GB available. The container was started 24 seconds ago. The engine is running.

STAGE 4: KUBERNETES DEPLOYMENT

WORKFLOW OVERVIEW:

Kubernetes is used to orchestrate the containerised Flask API at scale. The deployment.yml file defines a Deployment with 3 replicas (pods) of the travel-api container and a LoadBalancer Service to expose the deployment externally.

HOW IT WORKS:

1. The deployment.yml is applied using `kubectl apply`.
2. Kubernetes creates 3 identical pods from the travel-api Docker image, each running the Flask API independently.
3. A LoadBalancer service distributes incoming traffic across all 3 pods, ensuring high availability and horizontal scalability.
4. If any pod crashes, Kubernetes automatically restarts it (self-healing).

SCREENSHOT:

```

• (base) karthikdoguparthi@Lazys-Mac Module 7 % kubectl apply -f deployment.yml
deployment.apps/travel-api-deployment unchanged
service/travel-api-service unchanged
• (base) karthikdoguparthi@Lazys-Mac Module 7 % kubectl get pods
kubectl get service travel-api-service

NAME                                READY   STATUS    RESTARTS   AGE
travel-api-deployment-6c5cdb6b89-5cbvj  1/1     Running   0           25s
travel-api-deployment-6c5cdb6b89-l8c5x  1/1     Running   0           25s
travel-api-deployment-6c5cdb6b89-rvgr9  1/1     Running   0           25s
NAME                                TYPE          CLUSTER-IP      EXTERNAL-IP    PORT(S)          AGE
travel-api-service LoadBalancer  10.96.51.190    172.18.0.5     80:32491/TCP     25s
○ (base) karthikdoguparthi@Lazys-Mac Module 7 %

```

kubectl output showing 3 pods Running and LoadBalancer service

The terminal shows the following commands and their output:

```
kubectl apply -f deployment.yml  
→ deployment.apps/travel-api-deployment unchanged  
→ service/travel-api-service unchanged
```

```
kubectl get pods  
→ travel-api-deployment-6c5cdb6b89-5cbvj 1/1 Running 0 25s  
→ travel-api-deployment-6c5cdb6b89-l8c5x 1/1 Running 0 25s  
→ travel-api-deployment-6c5cdb6b89-rvgr9 1/1 Running 0 25s
```

```
kubectl get service travel-api-service  
→ NAME: travel-api-service  
→ TYPE: LoadBalancer  
→ CLUSTER-IP: 10.96.51.190  
→ EXTERNAL-IP: 172.18.0.5  
→ PORT(S): 80:32491/TCP  
→ AGE: 25s
```

All 3 pods are in Running state with 0 restarts, confirming a successful Kubernetes deployment. The LoadBalancer is exposing port 80 externally and routing to port 32491 internally.

STAGE 5: APACHE AIRFLOW SCHEDULING

WORKFLOW OVERVIEW:

Apache Airflow is used to automate the ML pipeline as a scheduled Directed Acyclic Graph (DAG). The DAG orchestrates the weekly workflow: loading raw data, preprocessing features, and retraining the regression model. This ensures the model stays up to date with fresh data without manual intervention.

HOW IT WORKS:

1. Airflow is launched in standalone mode using "airflow standalone".
2. The DAG file (travel_dag.py) is placed in the Airflow dags/ directory.
3. The DAG defines 3 tasks in sequence: load_data → preprocess → train_model.
4. Each task is a PythonOperator calling a Python function.
5. The DAG is scheduled to run @weekly, with catchup=False to avoid backfilling.
6. The Airflow scheduler picks up the DAG and makes it visible in the web UI.

SCREENSHOT:

```
standalone | Complete
(base) karthikdoguparthi@Lazys-Mac Module 7 % airflow standalone
/Users/karthikdoguparthi/opt/anaconda3/lib/python3.9/site-packages/airflow/config
uration.py:2261 FutureWarning: The 'executor' setting in [core] has the old d
efault value of 'SequentialExecutor'. This value has been changed to 'LocalExec
utor' in the running config, but please update your config.
[2026-05-06T18:42:35.436+0530] {providers_manager.py:953} INFO - The hook_class
'airflow.providers.standard.hooks.filesystem.FSHook' is not fully initialized
(UI widgets will be missing), because the 'flask_appbuilder' package is not ins
talled, however it is not required for Airflow components to work
[2026-05-06T18:42:35.438+0530] {providers_manager.py:953} INFO - The hook_class
'airflow.providers.standard.hooks.package_index.PackageIndexHook' is not fully
initialized (UI widgets will be missing), because the 'flask_appbuilder' packa
ge is not installed, however it is not required for Airflow components to work
standalone | Starting Airflow Standalone
standalone | Password for the admin user has been previously generated in /User
s/karthikdoguparthi/Downloads/AlmaBetter/Capstone/Module 7/airflow/simple_auth_
manager_passwords.json.generated. Not echoing it here.
standalone | Checking database is initialized
standalone | Database ready
scheduler | [2026-05-06T18:42:37.759+0530] {providers_manager.py:953} INFO - T
he hook_class 'airflow.providers.standard.hooks.filesystem.FSHook' is not fully
initialized (UI widgets will be missing), because the 'flask_appbuilder' packa
ge is not installed, however it is not required for Airflow components to work
scheduler | [2026-05-06T18:42:37.760+0530] {providers_manager.py:953} INFO - T
he hook_class 'airflow.providers.standard.hooks.package_index.PackageIndexHook'
is not fully initialized (UI widgets will be missing), because the 'flask_appb
uilder' package is not installed, however it is not required for Airflow compon
ents to work
api-server | [2026-05-06T18:42:37.797+0530] {providers_manager.py:953} INFO - T
```

Terminal: Airflow standalone startup

The terminal shows the output of "airflow standalone" being executed. Key log lines confirm: the Airflow database is initialised and ready, the admin password has been previously generated and stored in simple_auth_manager_passwords.json, the standalone mode is starting successfully, and both the scheduler and api-server components are initialising. The FutureWarning about executor settings is a non-critical configuration note.

Note: The Airflow Web UI screenshot (showing the DAG graph view with green task nodes) should be added from the Airflow UI at <http://localhost:8080> after logging in and triggering the "travel_retraining_pipeline" DAG manually.

DAG STRUCTURE:

load_data → preprocess → train_model

Schedule: @weekly

Tags: travel, mlops

STAGE 6: CI/CD PIPELINE JENKINS

WORKFLOW OVERVIEW:

A Jenkins CI/CD pipeline automates the build, test, and deployment cycle. Every time code is pushed to the GitHub repository, Jenkins automatically: clones the repository, installs dependencies, runs tests, builds the Docker image, pushes it to DockerHub, and deploys the updated image to Kubernetes.


HOW IT WORKS:

1. Jenkins is configured with a Pipeline job pointing to the GitHub repository.
2. The Jenkinsfile in the root of the repo defines 6 stages.
3. DockerHub credentials are stored securely in Jenkins as a credential binding.

4. On each push, the pipeline triggers automatically via a webhook or manual trigger.
5. The final stage runs `kubectl apply` to update the Kubernetes deployment with the new image.

PIPELINE STAGES:

- Stage 1: Clone Repository `git clone` from GitHub main branch
- Stage 2: Install Dependencies `pip install -r requirements.txt`
- Stage 3: Run Tests `pytest tests/`
- Stage 4: Build Docker Image `docker build -t travel-api .`
- Stage 5: Push to DockerHub `docker push username/travel-api:latest`
- Stage 6: Deploy to Kubernetes `kubectl apply -f deployment.yml`



Jenkins UI at <http://localhost:8080>.

STAGE 7: MLflow TRACKING


WORKFLOW OVERVIEW:

MLflow is used to track all machine learning experiments, log model parameters and metrics, and register the best-performing model in the MLflow Model Registry. This enables reproducibility, version control, and systematic comparison of model runs.

HOW IT WORKS:


1. The `mlflow_tracking.py` script sets up a named experiment ("Travel_Price_Prediction").
2. Each training run is wrapped in an `mlflow.start_run()` context.
3. Model hyperparameters (`n_estimators`, `learning_rate`, `max_depth`) are logged with `mlflow.log_param()`.
4. Evaluation metrics (R^2 , MAE, RMSE) are logged with `mlflow.log_metric()`.
5. The trained model is saved and registered in the Model Registry using `mlflow.sklearn.log_model()` with `registered_model_name="flight_price_model"`.
6. The MLflow UI is accessed at <http://127.0.0.1:5000> to view all runs and the registered model.

SCREENSHOT:



MLflow Experiments View: Travel_Price_Prediction

The MLflow Experiments page shows the "Travel_Price_Prediction" experiment in the left sidebar (selected with blue checkbox). The runs list shows 1 matching run named "amazing-toad-56" with status Complete (green), created 2 minutes ago, duration 3.5 seconds, source file mlflow_tracking.py, and the registered model "flight_price_model" linked to the run.



MLflow Model Registry: flight_price_model

The MLflow Model Registry shows the registered model named "flight_price_model" under the experiment "Travel_Price_Prediction". The model details confirm:

- Status: Ready (green indicator)
- Model ID: m-fbcc5d3fee274c0daaaaa94e68ac0cbc
- Source run: amazing-toad-56
- Source run ID: c4f6b6a96be14d3aa614847b962125c3
- Logged from: mlflow_tracking.py

Metrics logged:

- mae: 2.05 (extremely low mean absolute error)
- r2: 0.9999999999999999 ($R^2 \approx 1.0$ model explains virtually all price variance)

Parameters logged:

- max_depth: 7
- learning_rate: 0.2
- n_estimators: 100

SUMMARY

All MLOps stages have been successfully implemented:

1. REST API (Flask): Two prediction endpoints: /predict_price and /predict_gender — both returning HTTP 200
2. Streamlit App: Interactive dashboard with EDA, hotel recommendations, and live predictions (\$1224.30 flight price, gender: none)
3. Docker Deployment: "travel-container" built from "travel-api" image, running on port 5001:5001, 104.2MB memory
4. Kubernetes Deployment: 3 replica pods (all Running, 0 restarts) behind a LoadBalancer on EXTERNAL-IP 172.18.0.5
5. Airflow Scheduling: Standalone Airflow running with travel_retraining_pipeline DAG (load_data → preprocess → train_model, @weekly)
6. CI/CD Pipeline: Jenkins Jenkinsfile with 6 automated stages: clone → install → test → build → push → deploy
7. MLflow Tracking: Experiment "Travel_Price_Prediction", run "amazing-toad-56", $R^2 \approx 1.0$, MAE ≈ 2.05 , model registered as "flight_price_model"